

TSCI 5050
Introduction to Data Science
Fall 2018

CLASS DAYS and TIME: Alternate Tuesdays, 10 - Noon

CLASSROOM: LIB 2.028

COURSE DIRECTOR: Alex F. Bokov, PhD

OFFICE LOCATION and HOURS: Flexible, by appointment

EMAIL: bokov@uthscsa.edu

TELEPHONE: 210 562-4106

READ THIS DOCUMENT CAREFULLY – YOU ARE RESPONSIBLE FOR ITS CONTENTS

COURSE DESCRIPTION AND OBJECTIVES

This elective course is designed to train students to use the R programming languages to extract, prepare, and analyze data. This course is designed to be self-contained: statistical methods and theory relevant to analyzing large datasets will be covered with the computer-related course content providing tangible applications and motivating examples. In addition, the course will include organizational skill training and best practices needed to run a successful collaboration between researchers conducting patient oriented clinical research and researchers in computational fields.

Pre-requisites – None

Semester credit hours – 1.0

Learning Objectives:

- Aim: students will gain a basic working knowledge of programming languages, using R as a case study illustrating general concepts of algorithmic reasoning.
 - Students will translate a mathematical function that they understand into the form of executable code in the R language.
 - Given an R function that contains an error, students will use `debug()` and `browser()` along with reductionist reasoning to step through the function and locate the error.
- Aim: students will learn how to use the general linear model and the idea that statistics is a formalization of the scientific method to make sense of new statistical concepts.
 - Given a dataset and a verbal summary of the scientific goals of the study, students will fit a multivariable linear model using R, perform variable selection, and for each main effect or first-order interaction in the model write a sentence explaining its practical/clinical interpretation.
 - Given the default analytic output for each term in the above model, students will give the estimated effect-size, standard error, and whether or not that effect is statistically significant.

Students will use visual methods to diagnose potential problems with the above analysis and write a paragraph that can unambiguously communicate to a statistician what was already attempted and what problems require specialist assistance.

COURSE ORGANIZATION

The main teaching modalities used in this course include:

1. Didactic
2. Hands-on tutorials

Materials:

YOU WILL NEED THE FOLLOWING ONLINE ACCOUNTS (all are free):

- RStudio Cloud = <https://login.rstudio.cloud/register>
- GitHub = <https://github.com/join>
- Stack Overflow (recommended) = <http://stackoverflow.com/users/login>

The following link gives you step by step instructions for how to sign up for the two required accounts **and correctly link them to each other**-- https://github.com/bokov/TSCI-5050-Course-Notes/raw/master/TSCI5050_account_setup_instructions.pdf

Do not contact IMS for help with course-specific software. Please contact the course director instead, bokov@uthscsa.edu

You are encouraged bring a dataset from your own research that you are having trouble with, to use as part of the coursework. You will get a lot more out of the course that way.

Homework:

Homework is assigned on a topic *before* that topic is covered. That is deliberate, because it will help you absorb the material better after having already tried to do it on your own.

COMPUTER REQUIREMENTS

Students are required to have a laptop computer that can connect to and operate over a wireless network. All laptops will connect to The UTHSCSA network via the HSCwave broadcast wireless connection. Authentication for wireless use is based on The UTHSCSA domain username and password.

Verification of proper operation prior to the start of class is highly recommended. Assistance with networking is available thru the IMS Service Desk

- Telephone: 567-7777
- E-mail ims-servicedesk@uthscsa.edu)

Assistance is also available at the IMS Student Support Center (ALTC 106). But again, only bug them about standard UTHSCSA software, not the weird course-specific stuff listed above.

Reading Assignments – Reading assignments will be listed in the individual class sections of this syllabus.

ATTENDANCE

Attendance at scheduled classes and examinations is crucial to meeting course objectives. Therefore, regular attendance in class is expected of each student.

- Attendance is defined as being present within 15 minutes after the scheduled beginning of the class and until 15 minutes before the scheduled ending of the class.
- Excused absences may be granted by the Course Director in cases such as formal presentations at scientific meetings, illness, or personal emergency.

- Excused absences are considered on an individual basis and require electronic communication with the Course Director to request an excused absence. The e-mail request to the Course Director for consideration of an excused absence must provide details regarding the circumstances and specific dates.
- It is expected that students will provide *advanced notice* of absence for scheduled events.
- If a student has excessive unexcused absences in a given course, they will automatically receive a grade of *unsatisfactory* unless *makeup* has been approved by the Course Director.
- Makeup of absences (both excused and unexcused) is allowed at the discretion of the Course Director.
- Allowable unexcused absences will be determined by the credit hours of the course as follows:

Course Semester Credit Hours	Allowable Unexcused Absences
3.0	3
2.0	2
1.0	1

TEXTBOOKS

Strictly Optional:

- Reading materials will be provided. The following are good books to own, but not required for the class:
 - Dalgaard P. Introductory statistics with R. 2nd ed. New York: Springer; 2008.
 - ISBN: 978-0-387-79053-4
 - Description: Light and easy.
 - Burns P. S poetry. Lulu Com; 2012.
 - ISBN: 978-1-4710-4552-3
 - Description: Small but wise. Like Yoda.
 - Crawley MJ. The R book. Chichester, England?; Hoboken, N.J: Wiley; 2007.
 - ISBN: 978-0-470-51024-7
 - Description: Heavier, more detailed.
 - Pinheiro JC, Bates DM. Mixed-effects models in S and S-PLUS. New York: Springer; 2000.
 - ISBN: 978-0-387-98957-0
 - Description: Reality check for the over-confident.

GRADING POLICIES AND EXAMINATION PROCEDURES

Satisfactory/Unsatisfactory based on best 4 assignments and class participation via GitHub and RStudio Cloud.

Grading System

The grading will be conducted on a pass fail basis and both assignments need a Satisfactory in order to pass the course. S = Satisfactory U = Unsatisfactory

REQUESTS FOR ACCOMODATIONS FOR DISABILITIES

In accordance with policy 4.2.3, **Request for Accommodation Under the ADA and the ADA Amendments Act of 2008 (ADAAA)**, any student requesting accommodation must submit the appropriate request for accommodation under the American with Disabilities Act (ADA, form 100). To his/her appropriate Associate Dean of their School and a copy to the ADA Coordinator. Additional information may be obtained at <http://uthscsa.edu/eeo/request.asp>.

ACADEMIC INTEGRITY AND PROFESSIONALISM

Any student who commits an act of academic dishonesty is subject to discipline as prescribed by the UT System Rules and Regulations of the Board of Regents. Academic dishonesty includes, but is not limited to, cheating, plagiarism, collusion, the submission for credit of any work or materials that are attributable in whole or in part to another person, taking an exam for another person, signing attendance sheets for another student, and any act designed to give unfair advantage to a student or the attempt to commit such an act. Additional information may be obtained at

<http://catalog.uthscsa.edu/generalinformation/generalacademicpolicies/academicdishonestypolicy/>

TITLE IX AT UTHSCSA

Title IX Defined:

Title of the Education Amendments of 1972 is a federal law that prohibits sex discrimination in education. It reads “no person in the United States shall, on the basis of sex, be excluded from participation in, be denied the benefits of, or be subjected to discrimination under any education program or activity receiving Federal financial assistance.”

University of Texas Health Science Center San Antonio’s Commitment:

University of Texas Health Science Center San Antonio (UTHSCSA) is committed to maintaining a learning environment that is free from discriminatory conduct based on gender. As required by Title IX, UTHSCSA does not discriminate on the basis of sex in its education programs and activities, and it encourages any student, faculty, or staff member who thinks that he or she has been subjected to sex discrimination, sexual harassment (including sexual violence) or sexual misconduct to immediately report the incident to the Title IX Director.

In an emergency, victims of sexual abuse should call 911. For non-emergencies, they may contact UPD at 210-567-2800. Additional information may be obtained at <http://students.uthscsa.edu/titleix/>

EMAIL POLICY

All correspondence will be sent to the student using the student’s LiveMail address and CANVAS. All correspondence from the student to the course director should be sent to the course director’s e-mail as listed on the first page of this syllabus.

USE OF RECORDING DEVICES

Encouraged. Please send course director a copy!

ELECTRONIC DEVICES

Cell phones must be put on silent mode during all class meetings and exams. Computers and other devices should be used only for participating in classroom activities.

TENTATIVE CLASS SCHEDULE

**TSCI 5050
Introduction to Data Science
Fall 2018**

Week	Date	Module	Title/Instructor(s)
1	August 20th	Intro	Introduction to RStudio
2	September 3rd	R	GitHub, git, and bash
3	September 17th	R	Anatomy of a Programming Language, Part 1
4	October 1st	SQL	Anatomy of a Programming Language, Part 2
5	October 15th	Stats	Data Storage, Manipulation, and Presentation
6	October 29th	Stats	Stats Crash Course 1, (almost) everything is some form of regression
7	November 12th	Stats	Stats Crash Course 2, Real life is never tidy
8	November 26th	Wrap-Up	Stats Crash Course 3, TBA
9	December 10th	Wrap-Up	Special Topics, Final Questions

Week: 1
Date: August 20, 2019
Room: LIB 2.028
Instructor(s): Bokov
Topic: Introduction to RStudio
Learning Objectives and Competencies– Participants will be able to: <ol style="list-style-type: none"> 1. Understand structure and expectations of the course; introductions 2. Articulate what goals you might wish to achieve using biomedical informatics 3. Share any relevant research projects and experiences with the class 4. Understand how to find your way around your file system with RStudio 5. Develop a familiarity with labor-saving and error-reducing features of RStudio
Class Assignment: <ol style="list-style-type: none"> 1. Before and during first class: complete the instructions given in https://github.com/bokov/TSCI-5050-Course-Notes/raw/master/TSCI5050_account_setup_instructions.pdf 2. Problem set #1: discover intent of unfamiliar samples of R code 3. In one or two sentences, describe a data problem you face in your own work that you hope to solve computationally
Readings: http://blog.scottlowe.org/2015/01/14/non-programmer-git-intro/ Sections from free e-book <u>S-Poetry</u> indicated in the homework #1 handout.

Week: 2
Date: September 3, 2019
Room: LIB 2.028
Instructor(s): Bokov
Topic: GitHub, git, and bash
Learning Objectives and Competencies– Participants will be able to: <ol style="list-style-type: none"> 1. Configure GitHub for password-less login from RStudio Cloud 2. Be able to use git to clone, checkout, commit, and push. And explain what each of them does.
Class Assignment: <ol style="list-style-type: none"> 1. May opt to revise problem-set #1
Readings: TBA

Week: 3
Date: September 17, 2019
Room: LIB 2.028
Instructor(s): Bokov
Topic: Anatomy of a Programming Language, Part 1
Learning Objectives and Competencies– Participants will be able to: <ol style="list-style-type: none"> Understand the features that almost all programming languages have in common Be able to perform basic arithmetic, flow-control, and convenience/utility tasks in R Know where to go to find answers to R and what other resources exist online for improving your skills as you go along. Be able to turn a script into a function. Be able to create an empty function and write it from the “inside out” using the <code>browser()</code> command
Class Assignment: <ol style="list-style-type: none"> May opt to revise problem set #1 Finish customizing the code for uploading your individual dataset and declare your main analytic variables.
Readings: TBA

Week: 4
Date: October 1, 2019
Room: LIB 2.028
Instructor(s): Bokov
Topic: Anatomy of a Programming Language, Part 2
Learning Objectives and Competencies– Participants will be able to: <ol style="list-style-type: none"> Use the <code>lapply()</code> and <code>sapply()</code> functions instead of loops. Use the <code>%>%</code> operator to keep nested function calls readable. Demonstrate the use of all the subscripting operators. Use the <code>split()</code> command.
Class Assignment: <ol style="list-style-type: none"> May opt to further revise problem set #1 Customize your script to produce a data characterization table (not the data dictionary) and a visualization of your data. Problem set #2
Readings: TBA

Week: 5
Date: October 15, 2019
Room: LIB 2.028
Instructor(s): Bokov
Topic: Data Storage and Manipulation
Learning Objectives and Competencies– Participants will be able to: <ol style="list-style-type: none"> 1. Be able to explain the purpose of the main functions of the <code>dplyr</code> package. 2. Describe the main steps through which your data travels from file to report output.
Class Assignment: May opt to revise problem set #2
Readings: TBA

Week: 6
Date: October 29, 2019
Room: LIB 2.028
Instructor(s): Bokov
Topic: Stats Crash Course 1, (almost) Everything is Some Form of Regression
Learning Objectives and Competencies– Participants will be able to: <ol style="list-style-type: none"> 1. Be able to explain what a test statistic is, what a statistical distribution is, what they have to do with hypothesis tests, and what p-values actually are. 2. Be able to use visual analogies to explain the relationship between T-tests, ANOVA, and various types of regression. 3. Know the commands for fitting a simple linear regression model in R and interpret the most important parts of the output it yields. 4. Demonstrate basic skills for getting some benefit out of a quantitative journal article without necessarily understanding all the proofs.
Class Assignment: Problem set #3
Readings: (time permitting) Pick one from a TBA selection of peer-reviewed journal articles, critique use of statistics, and highlight the parts you don't understand

Week: 7
Date: November 12, 2019
Room: LIB 2.028
Instructor(s): Bokov
Topic: Stats Crash Course 2, Real Life Is Never Tidy
Learning Objectives and Competencies– Participants will be able to: <ol style="list-style-type: none"> 1. Integrate the material learned from previous units to understand what analysis-ready data should look like. 2. Be able to use residual plots and other diagnostic methods to validate your analysis, identify problems, and know what to do about them. 3. Be able to use contrasts to convert the default hypothesis tests of a regression model into ones that are relevant to your actual hypotheses. 4. Understand how to effectively collaborate with statisticians.
Class Assignment: <ol style="list-style-type: none"> 1. Perform diagnostic procedures learned in class on the analysis done in problem set #3 2. Attempt remedial measures where applicable
Readings: TBA

Week: 8
Date: November 26, 2019
Room: LIB 2.028
Instructor(s): Bokov
Topic: Stats Crash Course 3, Survival Analysis and Time Series
Learning Objectives and Competencies– Participants will be able to: <ol style="list-style-type: none"> 1. Be able to show examples other than mortality of where survival analysis is an appropriate approach 2. Know the commands for fitting a simple survival model in R and interpret the most important parts of the output it yields 3. Know how to perform diagnostics on survival models and what to do about them.
Class Assignment: If a time-to-event approach makes sense for the dataset you have been working on in previous homeworks, try this approach. Otherwise, use one of the example datasets that come with R. Prepare final questions to ask next class.
Readings: TBA

Week: 9
Date: December 10, 2019
Room: LIB 2.028
Instructor(s): Bokov
Topic: Wrap-Up
Learning Objectives and Competencies– Participants will be able to: 1. Question and answer session. 2. If time permits, special topics: dashboards, intermediate ggplot, R code that writes more R code, ...
Class Assignment: Have a great winter break and keep in touch!
Readings: None